Time Series Analysis for Temperature Prediction

DATS 6313

Liang Gao

George Washington University

10 December 2023

Content

1 Overview

2 Preprocessing the Data

3 Stationarity

4 Time Series Decomposition

5 Holt-Winter Method

6 Feature Selection & Regression

7 Base Models

8 SARIMA Model

9 Residual Analysis

10 Conclusion

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Overview

- The Jena Weather dataset was recorded every 10 minutes(2004 - 2020)
- Choose 3 years (2018 - 2020) of data, and average the data in each hour
- The dataset has 22 columns including " date" and 21 numerical variables(e.g. atmospheric pressure, Relative Humidity, Vapor pressure).
- The dependent variable is the temperature in Celsius

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Preprocessing the Data

- Use the Drift method to fill in the missing values
- Outliers fix: Average method for CO2 & Wind speed; Naive for max PAR

Wind Speed

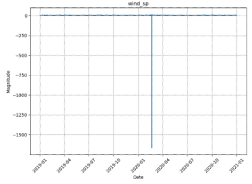

Max PAR

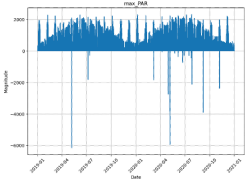

CO2

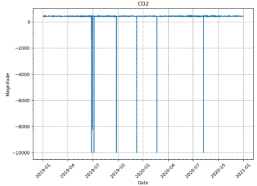

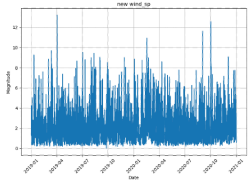

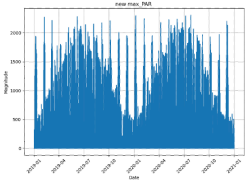

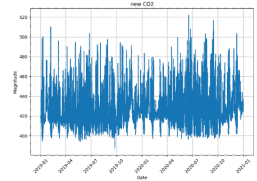

Preprocessing the Data

- Temperature over Time

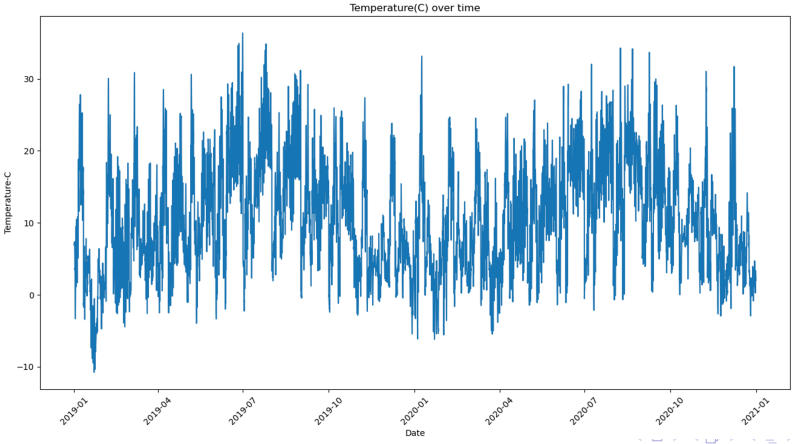

Preprocessing the Data

- ACF of Temperature

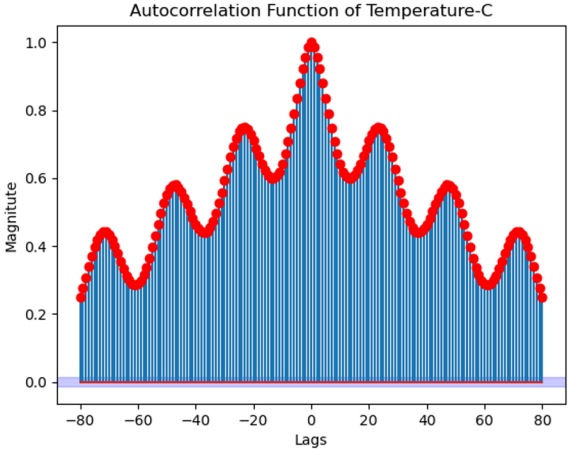

Preprocessing the Data

- The observation is 14,036 in the train set(80%) and 3,509 in the test set(20%)

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Stationarity

- The target variable passes the ADF test with a p-value of 0.00 but fails to pass the KPSS test with a p-value of 0.02
- The rolling mean and variance of temperature in Celsius, which stabilize once all samples are included
- The target variable dataset is weak-stationary

Stationarity

- Rolling mean & Variance of Temperature in Celsius

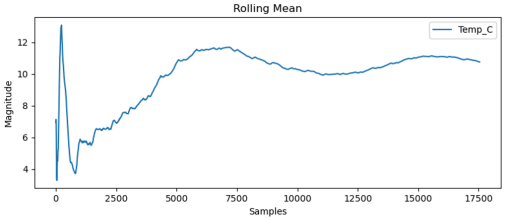

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Time Series Decomposition

- The strength of the trend is 94.37%, and the strength of the seasonality is 74.79%

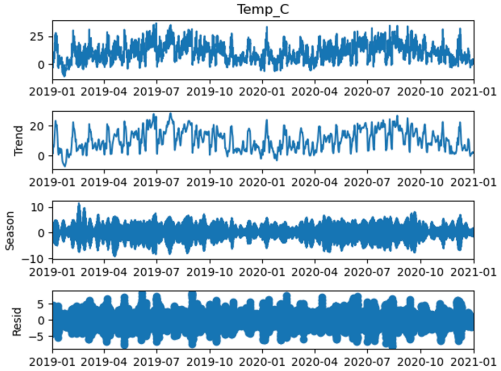

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Holt-Winter Method

- This method captures most seasonality but not the trend

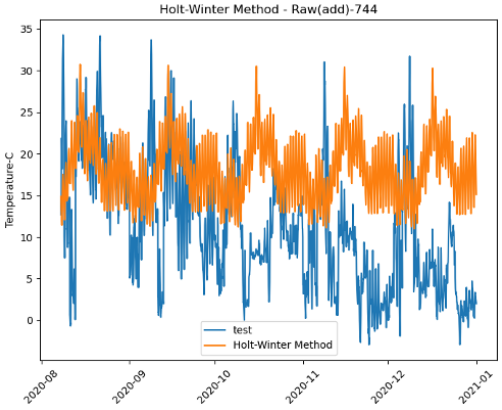

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Co-linearity check

- All the singular values are greater than 0, but the last few singular values are relatively small compared to the first largest one
- The condition number is 1,409,780.69 and is highly greater than 1,000
- Both results indicate severe co-linearity among some independent variables

Principal Component Analysis(PCA)

- The threshold for the PCA feature selection is a variance ratio of less than 0.95
- 7 features are chosen
- Adjusted R-squared - 0.982
- Mean of error - 0.005
- Variance of error - 0.017
- MSE - 0.017
- All the coefficients are statistically significant with p-values less than 0.05

Backwards Stepwise Regression

- Started with the model containing all independent variables, removed one predictor with the highest p-value at a time. 3 features were deleted

Remove	p-value	Adj_R2
\	\	1.00
PAR	0.79	1.00
Vapor_p_max	0.33	1.00
CO2	0.09	1.00

- Remove features with small coefficients(confidence interval for " rain time" is [-0.001, -0.000])

Backwards Stepwise Regression

- 8 features are chosen
- Adjusted R-squared: 1
- Mean of error: 0.001
- Variance of error & MSE: less than 0.00001
- All the coefficients are statistically significant with p-values less than 0.05
- Problem: The condition number of the regression model is 4.1e+03, which indicates strong multi-collinearity or other numerical problems

Variance Inflation Factor(VIF))

- The threshold for the VIF value is 10
- removed one predictor with the highest VIF value at a time(deleted 9 features)

remove	VIF	Adj_R2
		1.00
Vapor_p_max	14,403,743.53	1.00
H2O_conc	1,664,251.75	1.00
Vapor_p	18,405.61	1.00
PAR	790.68	1.00
air_density	304.34	1.00
Tlog	40.65	1.00
Temp_C_humi	24.89	0.97
wind_sp_max	24.69	0.97
SWDR	19.52	0.97

- Delete one insignificant feature & 6 features with small coefficients

Variance Inflation Factor(VIF)

- 3 features are chosen
- Adjusted R-squared: 0.971
- Mean & Variance of error: 0.031
- MSE: 0.030
- All the coefficients are statistically significant with p-values less than 0.05

Final Regression Model(VIF)

- Model derived from VIF has fewer features and no multi-collinearity problem
- Model performance

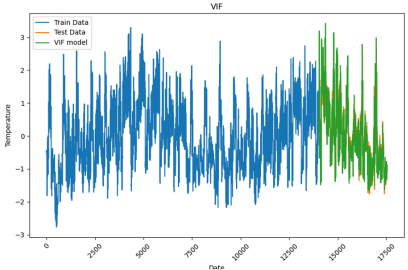

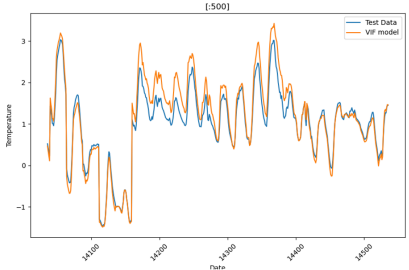

Final Regression Model(VIF)

- Hypothesis tests: F-test & T-test

T-test

	coef	std err	t	P>|t|	[0.025	0.975]
c0	-0.0063	0.001	-4.336	0.000	-0.009	-0.003
c1	-0.3764	0.003	-118.714	0.000	-0.383	-0.370
c2	0.2389	0.003	70.797	0.000	0.232	0.246
c3	0.7244	0.002	378.391	0.000	0.721	0.728

F-test

```
F-Test Results:
<F test: F=117357.3159157646, p=0.0, df_denom=1.4e+04, df_num=4>
```

Final Regression Model(VIF)

- Cross-validation

Subset	MSE	MeanRMSE	R-squared	Adj R-squared
1	0.05	0.22	0.96	0.96
2	0.04	0.19	0.97	0.97
3	0.02	0.15	0.97	0.97
4	0.04	0.19	0.97	0.97
5	0.03	0.16	0.97	0.97

- The consistency of the metrics across different subsets suggests that the model is stable and generalizes well to different subsets of the data

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Base Models

- Average, Niave, Drift, and SES

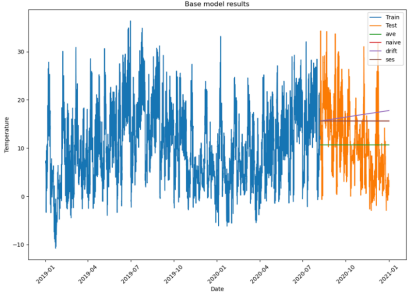

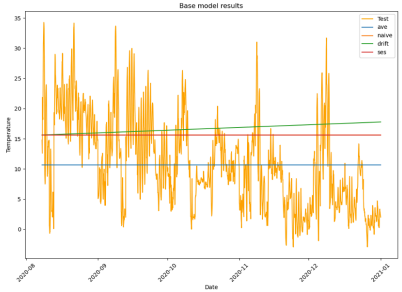

Base Models

Model	Mean	Variance	MSE
Average	0.39	55.47	55.62
Naive	-4.56	55.47	76.23
Dirft	-5.64	61.04	92.82
SES	-4.56	55.47	76.23

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

SARIMA Model

- GPAC & ACF/PACF of Raw Dataset

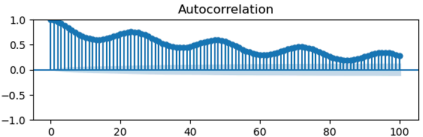

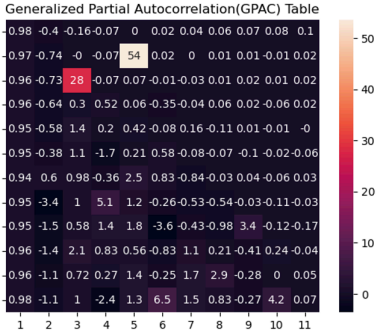

(1,0,1) (1,0,0,24)

- 1-step prediction & residual ACF

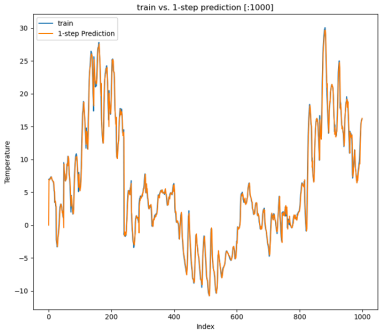

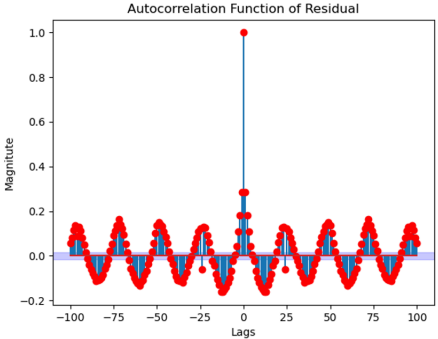

(1,0,1) (1,0,0,24)

- Residual ACF/PACF & GPAC

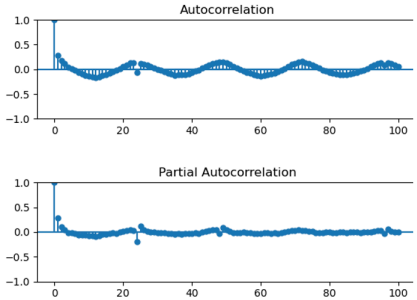

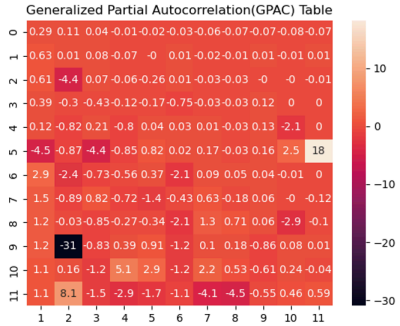

(1,0,3) (1,0,1,24)

- 1-step prediction & residual ACF

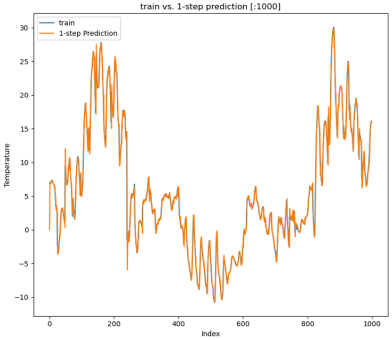

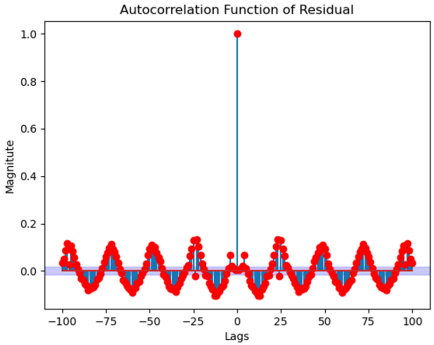

(1,0,3) (1,0,1,24)

- Residual ACF/PACF & GPAC

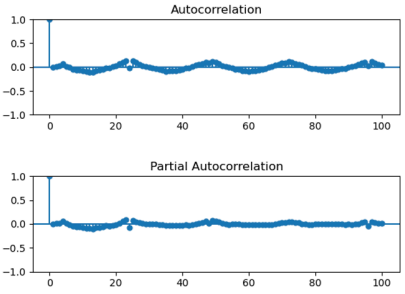

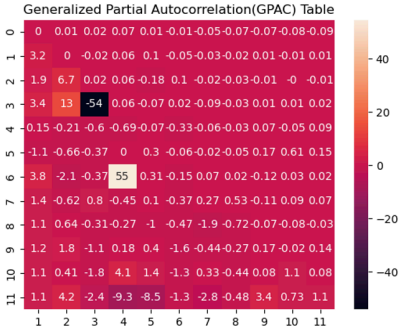

(1,0,3) (2,0,2,24)

- 1-step prediction & residual ACF

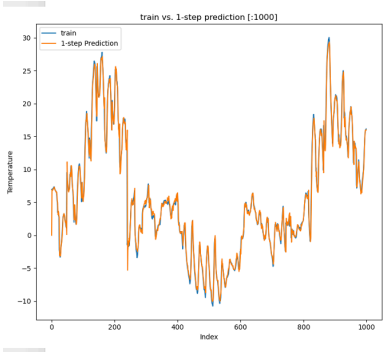

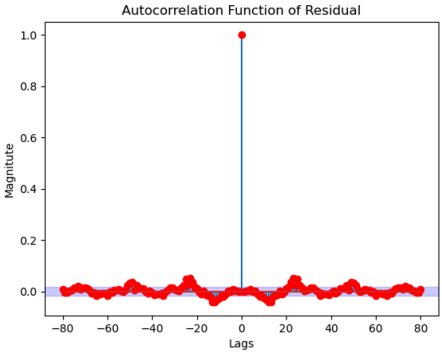

(1,0,3) (2,0,2,24)

- Residual ACF/PACF & GPAC

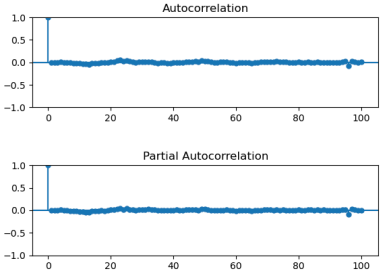

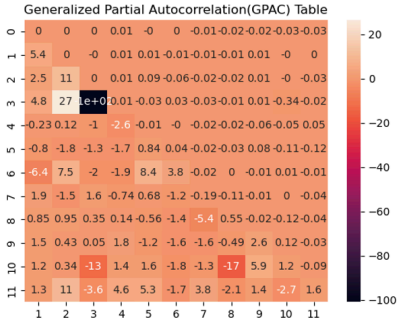

(1,0,3) (2,0,2,24)

- SARIMA model performance

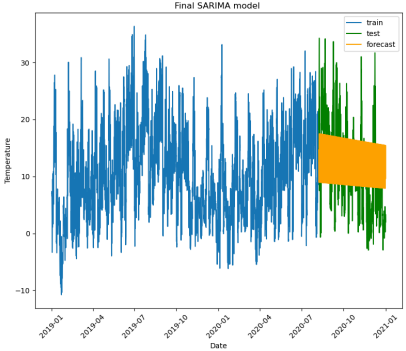

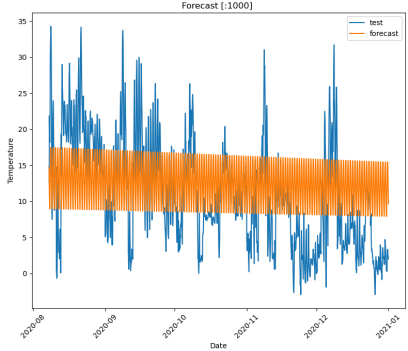

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Residual Analysis

- Box-Pierce test: Q $>$ Q*, fail the test
- Ljung-Box: p-values less than 0.05, fail the test
- Biased model: The estimated mean of the forecast error is -1.64
- Variance of the residual errors is 1.35 & Variance of forecast errors is 47.65
- Perform a zero-pole cancellation operation and there is no zero cancellation

List of Content

1 Overview
2 Preprocessing the Data
3 Stationarity
4 Time Series Decomposition
5 Holt-Winter Method
6 Feature Selection & Regression
7 Base Models
8 SARIMA Model
9 Residual Analysis
10 Conclusion

Conclusion

Method	Variance	Variance improvement (%)	MSE	MSE improvement (%)
Average	55.47	14.1 %	55.62	9.46%
Naive	55.47	14.1 %	76.23	33.94%
Dirft	61.04	21.94 %	92.82	45.74 %
SES	55.47	14.1%	76.23	33.94%
SARIMA	47.65	-	50.36	-